

Usage: memtier_benchmark [options]

A memcache/redis NoSQL traffic generator and performance benchmarking tool.

Connection and General Options:

-s, --server=ADDR	Server address (default: localhost)
-p, --port=PORT	Server port (default: 6379)
-S, --unix-socket=SOCKET	UNIX Domain socket name (default: none)
-P, --protocol=PROTOCOL	Protocol to use (default: redis). Other supported protocols are memcache_text, memcache_binary.
-a, --authenticate=CREDENTIALS	Authenticate using specified credentials. A simple password is used for memcache_text and Redis <= 5.x. <USER>:<PASSWORD> can be specified for memcache_binary or Redis 6.x or newer with ACL user support.
--tls	Enable SSL/TLS transport security
--cert=FILE	Use specified client certificate for TLS
--key=FILE	Use specified private key for TLS
--cacert=FILE	Use specified CA certs bundle for TLS
--tls-skip-verify	Skip verification of server certificate
--sni=STRING	Add an SNI header
-x, --run-count=NUMBER	Number of full-test iterations to perform
-D, --debug	Print debug output
--client-stats=FILE	Produce per-client stats file
--out-file=FILE	Name of output file (default: stdout)
--json-out-file=FILE	Name of JSON output file, if not set, will not print to json
--hdr-file-prefix=FILE	Prefix of HDR Latency Histogram output files, if not set, will not save latency histogram files
--show-config	Print detailed configuration before running
--hide-histogram	Don't print detailed latency histogram

--print-percentiles 50,99,99.9) --cluster-mode --help --version	Specify which percentiles info to print on the results table (by default prints percentiles: Run client in cluster mode Display this help Display version information
---	--

Test Options:

-n, --requests=NUMBER -c, --clients=NUMBER -t, --threads=NUMBER --test-time=SECS --ratio=RATIO --pipeline=NUMBER --reconnect-interval=NUM --multi-key-get=NUM --select-db=DB --distinct-client-seed --randomize	Number of total requests per client (default: 10000) use 'allkeys' to run on the entire key-range Number of clients per thread (default: 50) Number of threads (default: 4) Number of seconds to run the test Set:Get ratio (default: 1:10) Number of concurrent pipelined requests (default: 1) Number of requests after which re-connection is performed Enable multi-key get commands, up to NUM keys (default: 0) DB number to select, when testing a redis server Use a different random seed for each client random seed based on timestamp (default is constant value)
---	--

Arbitrary command:

--command=COMMAND --command-key-pattern=G --command-ratio --command-key-pattern	Specify a command to send in quotes. Each command that you specify is run with its ratio and key-pattern options. For example: --command="set __key__ 5" --command-ratio=2 To use a generated key or object, enter: __key__: Use key generated from Key Options. __data__: Use data generated from Object Options. The number of times the command is sent in sequence.(default: 1) Key pattern for the command (default: R): G for Gaussian distribution. R for uniform Random.
--	---

S for Sequential.

P for Parallel (Sequential were each client has a subset of the key-range).

Object Options:

- d --data-size=SIZE Object data size (default: 32)
- data-offset=OFFSET Actual size of value will be data-size + data-offset
- Will use SETRANGE / GETRANGE (default: 0)
- R --random-data Indicate that data should be randomized
- data-size-range=RANGE Use random-sized items in the specified range (min-max)
- data-size-list=LIST Use sizes from weight list (size1:weight1,..sizeN:weightN)
- data-size-pattern=R|S Use together with data-size-range
- when set to R, a random size from the defined data sizes will be used,
- when set to S, the defined data sizes will be evenly distributed across
- the key range, see --key-maximum (default R)
- expiry-range=RANGE Use random expiry values from the specified range

Imported Data Options:

- data-import=FILE Read object data from file
- data-verify Enable data verification when test is complete
- verify-only Only perform --data-verify, without any other test
- generate-keys Generate keys for imported objects
- no-expiry Ignore expiry information in imported data

Key Options:

- key-prefix=PREFIX Prefix for keys (default: "mementier-")
- key-minimum=NUMBER Key ID minimum value (default: 0)
- key-maximum=NUMBER Key ID maximum value (default: 10000000)
- PATTERN Set:Get pattern (default: R:R)
- G for Gaussian distribution.
- R for uniform Random.
- S for Sequential.
- P for Parallel (Sequential were each client has a subset of the key-range).

--key-stddev

The standard deviation used in the Gaussian distribution
(default is key range / 6)

--key-median

The median point used in the Gaussian distribution
(default is the center of the key range)

WAIT Options:

--wait-ratio=RATIO

Set:Wait ratio (default is no WAIT commands - 1:0)

--num-slaves=RANGE

WAIT for a random number of slaves in the specified range

--wait-timeout=RANGE

WAIT for a random number of milliseconds in the specified range (normal distribution with the center in the middle of the range)